

# **Welcome to COMP4300/8300 Parallel Systems**

**Prof. John Taylor**

**February 2025**

# Teaching staff, communication, website, and week 1 checklist

- Course Convenor & Lecturer: Prof. [John Taylor](#)
- Tutors:
  - Xianghao Wang
  - Yumin Li
- Communication with teaching staff handled via [Ed Discussion](#)
  - (do not send course-related e-mails to our personal ANU accounts please!)
- Course webpage at <https://comp.anu.edu.au/courses/comp4300>
  - (we will use Wattle only for quizzes, lecture recordings and marking)
- Week 1 checklist
  - Go to [this page](#) and follow the steps there
  - [Labs](#) start in week 2 Register before the end of week 1 via [MyTimeTable](#)
  - Account at NCI/Gadi needed ([account creation instructions](#))
  - Read carefully course policy on the appropriate usage of allocated computed time available ([available here](#))

# CSS

# CLASS REPRESENTATIVES

Class Student Representation is an important component of the teaching and learning quality assurance and quality improvement processes within the ANU College of Systems and Society (CSS).

Each semester, we put out a call for Class Representatives for all ANU College of Systems and Society (CSS) courses. Students can nominate themselves for one or more of the courses they are enrolled in.



Australian  
National  
University

# Roles and responsibilities:

The role of Student Representatives is to provide ongoing constructive feedback on behalf of the student cohort to Course Conveners and to Associate Directors (Education) for continuous improvements to the course.

- Act as the official liaison between your peers and convener.
- Be available and proactive in gathering feedback from your classmates.
- Attend regular meetings, and provide reports on course feedback to your course convener
- Close the feedback loop by reporting back to the class the outcomes of your meetings.

Note: Class representatives will need to be comfortable with their contact details being made available via Wattle to all students in the class.

For more information regarding roles and responsibilities, contact: ANUSA CSS representatives ([sa.cecc@anu.edu.au](mailto:sa.cecc@anu.edu.au)).



Image credit: he



Image credit: he

## Why become a class representative?

- **Ensure students have a voice** to their course convener, lecturer, tutors, and College.
- **Develop skills sought by employers**, including interpersonal, dispute resolution, leadership and communication skills.
- **Become empowered.** Play an active role in determining the direction of your education.
- **Become more aware of issues influencing your University** and current issues in higher education.
- **Course design and delivery.** Help shape the delivery of your current courses, as well as future improvements for following years.

**Want to be a class representative?  
Nominate today!**

Please nominate yourself to your course convener by end of Week 2



# BECOME A SCIPALS MENTOR!

APPLICATIONS CLOSE 27TH FEBRUARY (5PM)




Are you a STEM student in your second year or beyond? Want to boost your ANU+ volunteering hours while making a meaningful impact? Interested in guiding and connecting with new STEM students? SciPals is an ANU+ accredited mentoring program designed to welcome and integrate new STEM students into the broader ANU science community. As a mentor, you'll be matched with mentee(s) who can greatly benefit from your insights and experience. We encourage anyone eager to share advice, develop mentoring skills, and meet new people. No prior mentoring experience is needed—just a willingness to share the tips and strategies that helped you navigate your STEM journey at ANU.

# Tech.Launcher\_

Launch.idea, Launch.start\_up, Launch.career\_

## ➤ HPC Cluster Design and Deployment



ST. LOUIS NOV 16-21

PROGRAMEXHIBITSSTUDENTS  
SCINET  
MEDIATTEND

Q


Home > Students@SC > Student Cluster Competition

## Student Cluster Competition

With sponsorship from hardware and software vendor partners, competing student teams design and build small clusters, learn scientific applications, and apply optimization techniques for their chosen architectures in a non-stop, 48-hour challenge.

### 48-HOUR HPC-A-PALOOZA

**Student Cluster Competition Schedule**  
Monday–Wednesday, November 17–19, 2025



# Course Assessment

- The total workload for this (6 units) course is around 150 hours (12 hours per semester week). For S1 2025, the following assessment scheme:
  - **Assignments – 50% of final mark**
    - There will be two assignments, each worth 25% of the final mark (50% in total)
    - These assignments will contain a significant parallel programming component plus a report where you will have to answer practical and theoretical questions relevant to the course content – start early!
  - **Quizzes – 10% of final mark**
    - Remote/electronic multiple-choice exam in Wattle
  - **Final Exam – 40% of final mark**
    - This will be a 3 hour written exam held at the end of the semester during the normal examination period
    - This exam is worth 40% of the final course mark



# Software and minimum IT requirements (I)

- You will need satisfactory Internet connection to be able to access to:
  - [course website](#) (lecture slides etc.)
  - [GitLab](#) (Labs; Assignments submission)
    - Some familiarity with Git/GitLab (clone, commit, push, pull, etc.) assumed
    - You will need a private/public key installed in your GitLab account to be able to work with Git repos ([instructions here](#))
    - Click [here](#) for a set of best practices and instructions on the workflow to be used for the assignments
  - [Wattle](#) (quizzes, marking, and lecture recordings)
  - [Ed Discussion](#)
  - NCI Gadi [documentation](#)

## Software and minimum IT requirements (II)

- You will also need:
  - An SSH client to access `gadi.nci.org.au` with your NCI account (if it gets as far as asking for a username/password outgoing SSH is not being blocked)
  - Later, you will need access to a GPU server (`stugpu2.anu.edu.au` via `partch.anu.edu.au`)
  - For code development, you may use a remote terminal editor (advanced users) or a graphical editor, in particular VSCode with the Remote Development extension is highly recommended for beginners. Click [here](#) for instructions.
  - For Windows users, [Moba XTerm](#) highly recommended to access Gadi

## Course recommended texts + extra materials

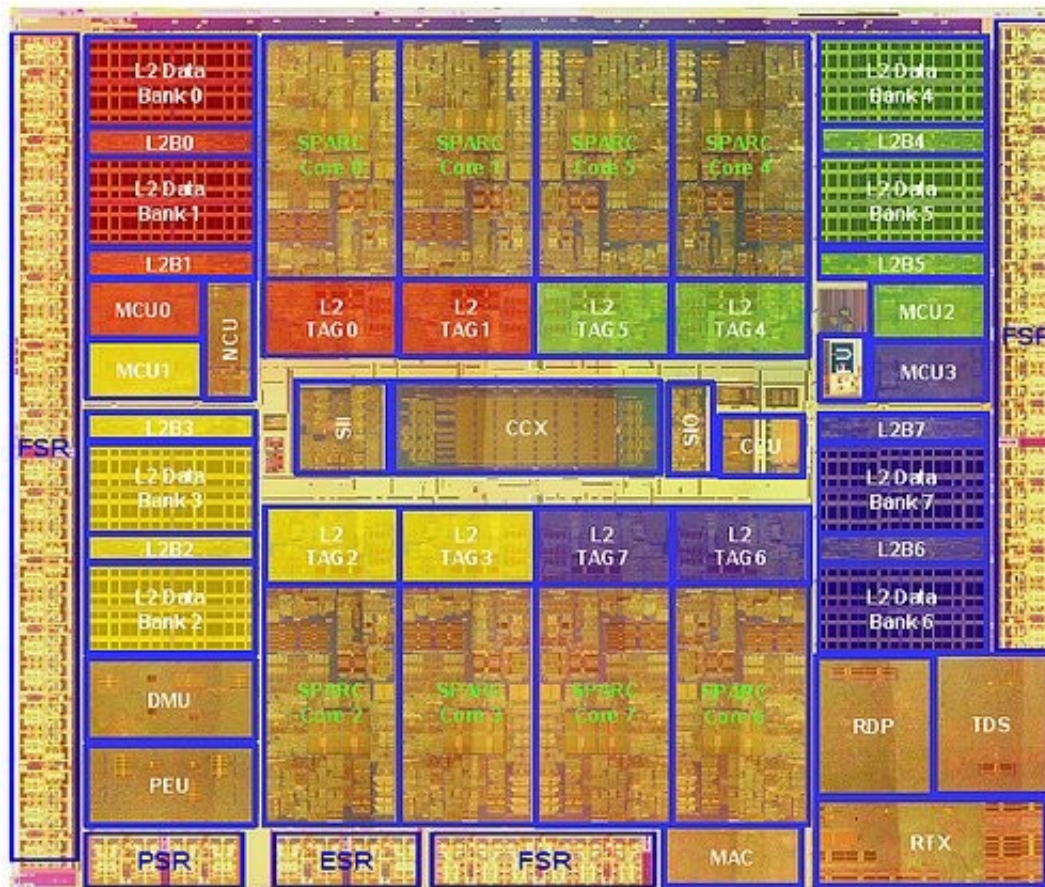
- The course does not strictly follow any particular textbook
- However, most of the concepts treated in the course are covered in the materials available at [resources section of the web page](#)

# Health Warning

- it's a 4000/8000-level course, it's supposed to:
  - be more challenging than a 3000-level course
  - you will be exposed to bleeding edge technologies
  - be less well structured
  - have a greater expectation for your self-directed understanding/independence
  - have more student participation
  - and why not ... be fun as well!
- it assumes you have done some background in concurrency (e.g. COMP2310); also in computer organization
- **COMP4300 requires strong programming skills – in C!**
- attendance at lectures/labs strongly recommended (even though not assessed)

# Today's lecture main take-away message

To extract all the computational performance from today's (even commodity) microprocessors there is no alternative but to **EXPLICITLY** manage parallelism in our programs!



UltraSPARC T2  
(2007)  
(Niagara-2)  
multicore  
chip layout

(courtesy of T.  
Okazaki, Flickr)







# The Cores Main Functions

## Instruction Processing

- Fetches instructions from memory
- Decodes instructions
- Executes instructions
- Manages instruction pipeline
- Handles branch prediction

## Arithmetic and Logic

- Performs mathematical calculations
- Handles logical operations
- Processes floating-point operations
- Executes integer operations

## Register Management

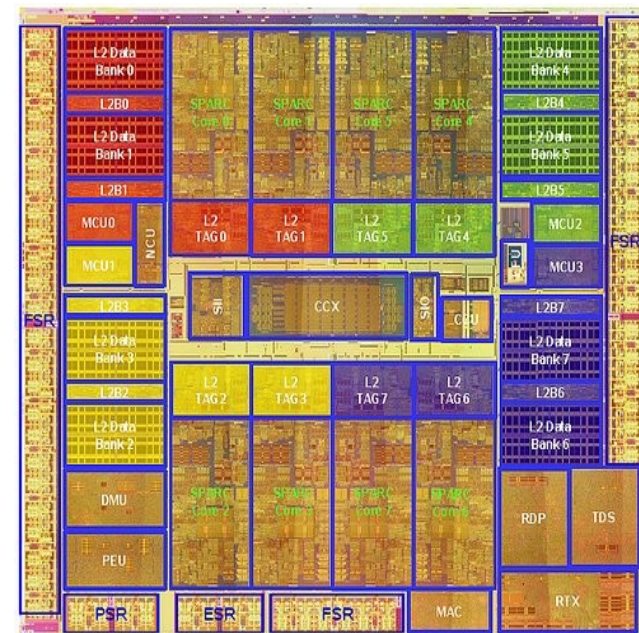
- Maintains program counters
- Manages general-purpose registers
- Handles special registers
- Controls status registers

## Control Flow

- Manages program execution flow
- Handles conditional operations
- Processes jump and branch instructions
- Deals with interrupts and exceptions

**This is a multi-core processor where each core can:**

- Work independently
- Execute different threads
- Share resources via the cache system (L2 cache)
- Communicate through the CCX (Cache Coherency Crossbar)



# Lecture Overview

- parallel computing concepts and scales
- sample application area: computational science and engineering (CSE)
- the role of Moore's Law and Dennard's Scaling Law in the exponential growth of computing performance
- end of Dennard's scaling Law (mid 2000s) and multicore revolution
- the Top 500 supercomputers and challenges for the future of computing
- why parallel programming is hard

# Parallel Computing: Concept and Rationale

## The idea:

Split computation into tasks that can be executed simultaneously on different processors

## Motivation:

- Speed, **Speed**, **Speed** . . . at a cost-effective price
  - if we didn't want it to go faster we wouldn't bother with the hassles of parallel programming!
- reduce the time to solution to acceptable levels (e.g., under real-time constraints)
  - no point in taking 1 week to predict tomorrow's weather!
  - simulations that take months are NOT useful in a design environment
- tackle larger-scale (i.e., with higher computational demands) problems
- keep power consumption and heat dissipation under control

# Parallelization

Split program up and run parts simultaneously on different processors

- on  $p$  computers, the time to solution should (ideally!) be reduced by  $\frac{1}{p}$
- **parallel programming**: the art of writing the parallel code
- **parallel computer**: the hardware on which we run our parallel code
  - *This course will discuss both!*
- Beyond raw compute performance, other motivations may include
  - enabling more accurate simulations in the same time (finer grids)
  - providing access to huge aggregate memories
  - providing more and/or better input/output capacity
  - hiding latency

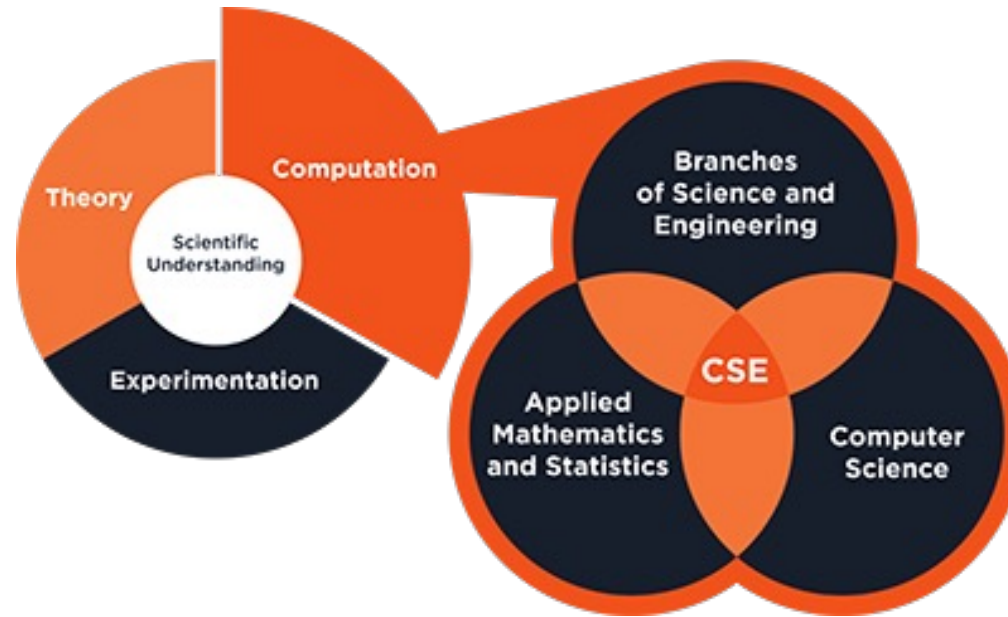


# Scales of Parallelism

- **within a CPU/core**: pipelined instruction execution, multiple instruction issue (superscalar), other forms of instruction level parallelism, SIMD units\*
  - **within a chip**: multiple cores\*, hardware multithreading\*, accelerator units\* (with multiple cores), transactional memory\*
  - **within a node**: multiple sockets\* (CPU chips), interleaved memory access (multiple DRAM chips), disk block striping / RAID (multiple disks)
  - **within a SAN** (system area network): multiple nodes\* (clusters, typical supercomputers), parallel filesystems
  - **within the internet**: grid/cloud computing\*
- \*requires significant parallel programming effort
- What programming paradigms are typically applied to each feature?

# Application Area

## Computational Science & Engineering (CSE)



### ➤ *Third pillar of scientific discovery*

- Integrates applied mathematics, computer science, and branches of science/engineering in a single discipline (e.g., computational geophysics)
- Leverages computational models, algorithms, data, software and HPC (i.e. **supercomputers**) to tackle grand-challenges in science and engineering

# Application Area - R&D in CSE


- Objective: improve the state-of-the-art in **computational models**, **algorithms**, and **software** to push the boundaries of what is currently achievable in CSE
- **Strong potential**: simulate out of reach problems, more precise predictive CSE, improved scientific knowledge, revolutionize decision-making across science, technology and society

## Main research areas

- Mathematical modelling
- Numerical methods (**discretization**, **solvers**)
- Data assimilation (e.g., **machine learning**)
- **HPC** (parallel software/hardware innovations)

## Application areas (examples)

- Geophysics
- Nuclear fusion
- Aeronautics
- Personalized medicine (brain/heart)
- Nanoscience, Smart manufacturing, (large) Etc.



# Application Area - Synergy among CSE and HPC

- We already find ourselves in the **Exascale** era ( $O(10^{18})$  FLOPs/s peak )
- **El Capitan**: Exascale supercomputer (Lawrence Livermore US National Labs) ( $\sim 11$ M cores, 1.7EFLOPs/s, ranked #1 Nov. 2024 [TOP500](#) list)




- Performance boost mostly based on adding hardware parallelism (e.g., higher #cores/CPU) and heterogeneous hardware (CPUs, GPUs, ...)
- To exploit such vast concurrency is a **formidable task** for CSE (breakthroughs in scalable algorithms and software innovations)

## Application Area

# The Fourth Paradigm: Data-intensive Scientific Discovery

- Data intensive science is considered a fourth paradigm of scientific discovery
- Book of this name published in 2009 focusing on the analysis of large data sets
- Machine learning/AI is now becoming an essential part of scientific research
- Big data is now the entire internet
- Parallel computing essential for LLM training





# Moore's Law & Dennard Scaling

Two “laws” underpin exponential performance increase of microprocessors



**Moore's Law**  
'Transistor density will double  
approximately every two years.'




Fig. 2 Number of components per integrated function for minimum cost per component extrapolated vs time.



**Dennard Scaling**  
'As MOSFET features shrink,  
switching time *and* power  
consumption will fall  
proportionately'




Fig. 1. Illustration of device scaling principles with  $k = 5$ . (a) Conventional commercially available device structure. (b) Scaled-down device structure.


# Moore's law

*The number of transistors in a chip doubles every 24 months*

Does this automatically imply that performance (per chip) doubles every 24 months?


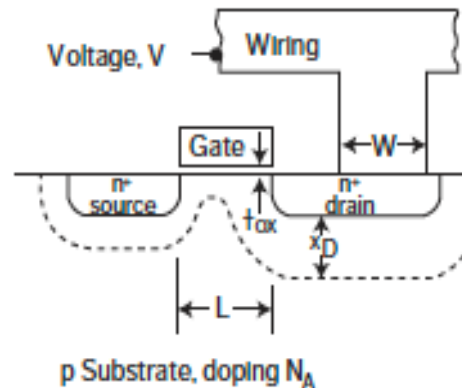
No! Moore's law is not a performance law! However, there has been a strong correlation between the # of transistors and performance across time...

*How is that possible?*



# Dennard's scaling law (until early 2000s)

Device or circuit parameter	Scaling factor
Device dimension $t_{ox}$ , $L$ , $W$	$1/\kappa$
Doping concentration $N_A$	$\kappa$
Voltage $V$	$1/\kappa$
Current $I$	$1/\kappa$
Capacitance $\epsilon A/t$	$1/\kappa$
Delay time/circuit $VC/I$	$1/\kappa$
Power dissipation/circuit $VI$	$1/\kappa^2$
Power density $VI/A$	1



**Figure 1.** Traditional MOSFET scaling as described by Robert Dennard. **Figure 2.** Minimum feature size scaling trend for Intel logic technologies.

The (dynamic) power consumption of a chip can be modelled as:

$$P = QfCV^2$$

where  $Q$  # of transistors,  $f$  frequency,  $C$  capacitance, and  $V$  voltage supply

# Dennard's scaling law (until early 2000s)

According to Dennard's law, if we scale feature size down by a factor of  $\frac{1}{K}$ , we can scale up frequency by  $K$ , and scale down the capacitance and voltage by  $\frac{1}{K}$ , resulting in a (reduced) power consumption of (assuming  $Q_K = Q_0$ ):

$$P_0 = Q_0 f_0 C_0 V_0^2 \rightarrow P_K = Q_K f_K C_K V_K^2 = Q_0 (K f_0) \left( \frac{1}{K} C_0 \right) \left( \frac{1}{K} V_0 \right)^2 = \left( \frac{1}{K^2} \right) P_0$$

What about if we allow ourselves to keep  $P_K$  constant? How many transistors  $Q_K$  can we fit on the same chip? The answer is:


$$Q_K = K^2 Q_0$$

As long as we keep scaling feature size down by  $\frac{1}{K}$ , we can fit  $K^2$  more transistors on the same chip, increase their frequency by  $K$ , and use the same power as before!

# Until early 2000s – The uniprocessor era

Performance of single instruction stream microprocessors increased exponentially  
(at a pace of  $\approx 40\%$  every 24 months)

- Exponential increase of clock rate
- Increase in Instruction Level Parallelism (ILP)
  - Pipelining: from 5 up to 31 (Prescott) stages
  - Superscalarity: from  $< 1$  instruction/clock to 4+ instructions/clock
- Many transistors for additional
  - optimizations
  - Increase in cache sizes
  - Prefetching
  - Sophisticated branch prediction logic



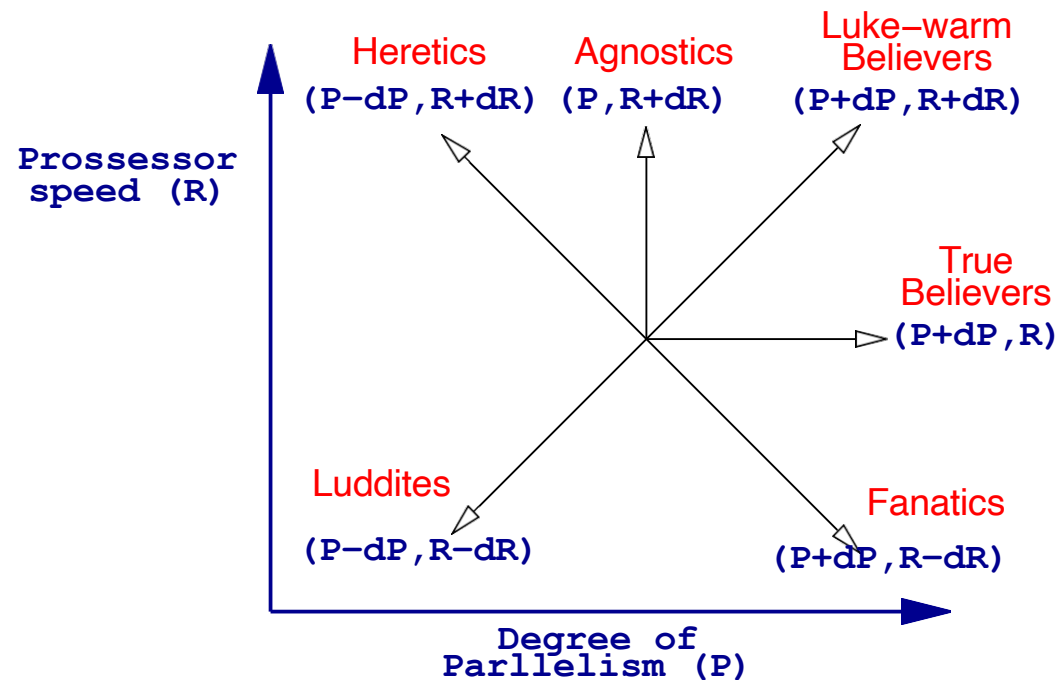


# Moore's Law and Dennard Scaling Undermine Parallel Computing

- Parallel computing looked promising in the 90's, but many companies failed due to the 'free lunch' from the combination of [Moore's Law](#) and [Dennard Scaling](#)
  - Why parallelize my codes? Just wait 2 years, and processors will be 1.4x faster!

*On several recent occasions, I have been asked whether parallel computing will soon be relegated to the trash heap . . . [Ken Kennedy, CRPC Director, 1994](#)*

- demography of Parallel Computing (mid 90's, origin unknown)



# The end of Dennard scaling and the uniprocessor era (2002-2004)


With feature size below  $\approx 100\text{nm}$  (nowadays  $< 5\text{nm}$ ), we have that:

$$P = QfCV^2 + VI_{\text{leakage}}$$

(Note: for “large enough” feature size, the term  $VI_{\text{leakage}}$  is negligible)

- Unfortunately,  $I_{\text{leakage}}$  grows exponentially with downscaled  $V$  as we decrease feature size by  $1/\kappa$ . Thus, the term  $VI_{\text{leakage}}$  blows up and dominates power consumption!
- To keep power under control, large number of transistors are switched off (dark silicon effect), operated at lower frequencies (dim silicon effect) or organized in different ways

# The end of Dennard scaling and the uniprocessor era




The free lunch is over!

- Clock rate capped ( $< 4$  GHz) by power constraints
- ILP (superscalarity and pipelining) reaches its limits
- Power usage saturated (becomes a major concern)

# The multicore era

The only way to sustain exponential growth in computational performance is  
by efficiently exploiting parallel computers



- This is very different from the golden years of ILP where hardware architects did all the work for us
- Programmers are now forced to bear the burden of finding and exploiting parallelism
- This is also an exciting era of opportunities for computational scientists: new algorithms and efficient implementations make a difference on what is achievable in computing

# Where are we now?

We find ourselves in a chip development era characterized by:

- Transistors galore (this cannot be sustained forever, though, due to physical limits)
- Severe power limitations (maximize transistors utility no longer possible)
- Customization versus generalization

We are now seeing:

- (customized) accelerators, generally manycore with low clock frequency
  - e.g. Graphics Processing Units (GPUs), customized for fast numerical calculations
- ‘dark silicon’: where we need to turn off parts of chip to reduce power
- hardware-software codesign: speed via specialization




# The Top 500 Most Powerful Computers: November 2024


The [TOP500](#) list provides an interesting view of these trends (click [here](#) for full list)

Rank	System	Cores	Rmax (PFlop/s)	Rpeak (PFlop/s)	Power (kW)
1	<b>El Capitan</b> - HPE Cray EX255a, AMD 4th Gen EPYC 24C 1.8GHz, AMD Instinct MI300A, Slingshot-11, TOSS, HPE DOE/NNSA/LLNL United States	11,039,616	1,742.00	2,746.38	29,581
2	<b>Frontier</b> - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE Cray OS, HPE DOE/SC/Oak Ridge National Laboratory United States	9,066,176	1,353.00	2,055.72	24,607
3	<b>Aurora</b> - HPE Cray EX - Intel Exascale Compute Blade, Xeon CPU Max 9470 52C 2.4GHz, Intel Data Center GPU Max, Slingshot-11, Intel DOE/SC/Argonne National Laboratory United States	9,264,128	1,012.00	1,980.01	38,698
4	<b>Eagle</b> - Microsoft NDv5, Xeon Platinum 8480C 48C 2GHz, NVIDIA H100, NVIDIA Infiniband NDR, Microsoft Azure United States	2,073,600	561.20	846.84	


# Top500: Performance Trends



# The Top500: Multicore Emergence



[\(http://www.top500.org/blog/slides-highlights-of-the-45th-top500-list/\)](http://www.top500.org/blog/slides-highlights-of-the-45th-top500-list/)



# Exascale and Beyond: Challenges and Opportunities

Level	Characteristic	Challenges/Opportunities
Among nodes	sheer #nodes, although different balance among #nodes and performance/node in top-ranked supercomputers (e.g., <a href="#">Frontier</a> has “only” 9.4K nodes but <a href="#">Fugaku</a> has 159K nodes)	<ul style="list-style-type: none"><li>➤ programming languages/environment</li><li>➤ fault tolerance</li></ul>
Within a node	High heterogeneity (e.g., <a href="#">Frontier</a> , <a href="#">LUMI</a> and many more use CPUs and GPUs	<ul style="list-style-type: none"><li>➤ what to use when</li><li>➤ co-location of data with the unit processing it</li></ul>
Within a chip	energy minimization (e.g. processor cores already have dynamic frequency and voltage scaling)	<ul style="list-style-type: none"><li>➤ minimize data size and movement (e.g., use just enough precision)</li><li>➤ specialized cores</li></ul>

# Why Parallel Programming is Hard

- Writing (correct and efficient) parallel programs is hard!
  - hard to expose enough parallelism; hard to debug!
- Getting (close to ideal) speedup is hard! Overheads include:
  - idling (e.g., caused by load unbalance, synchronization, serial sections, etc.)
  - redundant/extra operations when splitting a computation into tasks
  - communication time among processes
- Amdahl's Law: Let  $f$  the fraction of a computation that cannot be split into parallel tasks. Then, max speedup achievable for arbitrary large  $p$  (processors) is  $\frac{1}{f}$  !!!  
Let  $t_s$  and  $t_p$  denote serial and parallel exec times
$$t_p = f t_s + \frac{(1-f)t_s}{p}$$
$$S_p = \frac{t_s}{t_p} = \frac{p}{pf + (1-f)} \quad (\text{Speedup})$$
  - $\lim_{p \rightarrow \infty} S_p = \frac{1}{f}$
  - e.g., if  $f = 0.05$ , then  $\frac{1}{f} = 20$
- counterargument (Gustafson's Law):  $1 - f$  is not fixed, but increases with the data/problem size  $N$



## Recommended further reading

The following two articles provide a crystal clear view of past, present, and future:

- Computing Performance: Game Over or Next Level? Computer, 44(1), 2011. Publisher: IEEE.
- Time Moore: Exploiting Moore's Law From The Perspective of Time. IEEE Solid-State Circuits Magazine, 11 (1), 2019. Publisher: IEEE.

Only for intrepid readers:

- Computing Performance: Game Over or Next Level?, Full report from US National Research Council, National Academies Press, 2011. Bonus: reprints for Moore's and Dennard's Laws seminal papers.
- AI for next generation computing: Emerging trends and future directions *Internet of Things*, Elsevier, August 2022

## Additional references related to this lecture

- R. Dennard et al., “Design of Ion- Implanted MOSFETs with Very Small Physical Dimensions”. IEEE J. Solid State Circuits, vol. 9, no. 5, 1974, pp. 256–268.
- V. Agarwal, M.S. Hrishikesh, S.W. Keckler, and D.A. Burger, “Clock Rate versus IPC: The End of the Road for Conventional Microarchitectures”. In Proceedings of the 27th Annual International Symposium on Computer Architecture, 2000, pp. 248-259.
- M. T. Bohr and I. A. Young, ”CMOS Scaling Trends and Beyond”, in IEEE Micro, vol. 37, no. 6, pp. 20-29, November/December 2017, doi: 10.1109/MM.2017.4241347.
- M. B. Taylor, ”A Landscape of the New Dark Silicon Design Regime”, in IEEE Micro, vol. 33, no. 5, pp. 8-19, Sept.-Oct. 2013, doi: 10.1109/MM.2013.90.
- The Free Lunch Is Over. A Fundamental Turn Toward Concurrency in Software. Herb Sutter.  
<http://www.gotw.ca/publications/concurrency-ddj.htm>